

Homework 5 CART and Logistic Regression

Anushka Tak

Chandan Manjunath

Problem 1

```
setwd("C:/Users/pc/Desktop/Spring2019/DataMining/homework5")
cars<- read.csv("ToyotaCorolla.csv")
#head(cars)
dim(cars)

## [1] 1443 39
```

creating dummies for fuel type

```
library(fastDummies)
cars <- fastDummies::dummy_cols(cars, select_columns = "Fuel_Type")
#creating dummies for color
cars <- fastDummies::dummy_cols(cars, select_columns = "Color")
#head(cars)
```

splitting dataset into training, validation and test portions

```
sample_train<- sample(seq_len(nrow(cars)), size = floor(0.50*nrow(cars)))
sample_valid<- sample(seq_len(nrow(cars)), size = floor(0.30*nrow(cars)))
sample_test <- sample(seq_len(nrow(cars)), size = floor(0.20*nrow(cars)))

train      <- cars[sample_train, ]
validation<- cars[sample_valid, ]
test       <- cars[sample_test, ]

library(rpart)
library(rpart.plot)

dtm <- rpart(Price ~
Age_08_04+KM+Fuel_Type_Petrol+Fuel_Type_Diesel+Fuel_Type_CNG+Fuel_Type_HP+
Automatic+Doors+Quarterly_Tax+Mfr_Guarantee+
Guarantee_Period+Airco+Automatic_airco+CD_Player+
Powered_Windows+Sport_Model+Tow_Bar , method = "anova", data =
train)

dtm

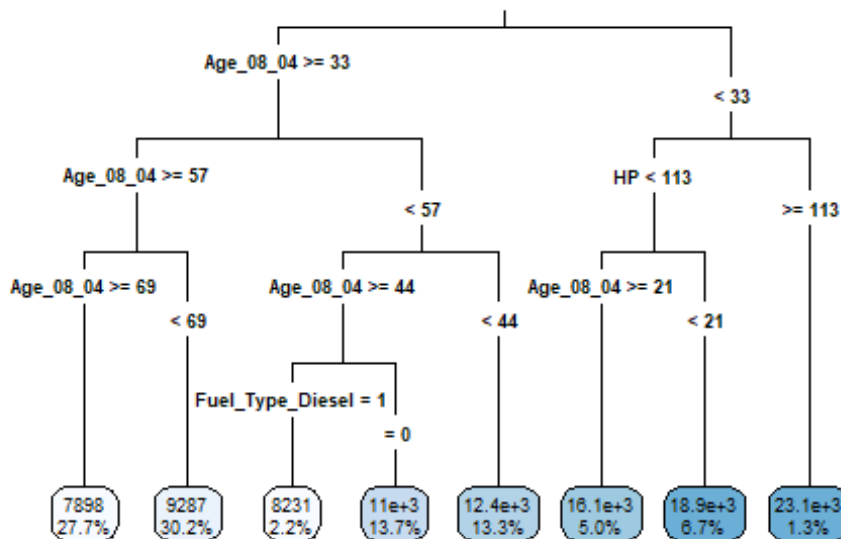
## n=716 (5 observations deleted due to missingness)
##
```

```

## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 716 9357383000 10689.160
##    2) Age_08_04>=32.5 623 2463283000 9563.592
##      4) Age_08_04>=56.5 414 671615400 8622.594
##        8) Age_08_04>=68.5 198 199775200 7897.672 *
##        9) Age_08_04< 68.5 216 272408200 9287.106 *
##      5) Age_08_04< 56.5 209 698920400 11427.580
##        10) Age_08_04>=43.5 114 360821100 10622.010
##          20) Fuel_Type_Diesel>=0.5 16 49339710 8230.938 *
##          21) Fuel_Type_Diesel< 0.5 98 205071100 11012.390 *
##        11) Age_08_04< 43.5 95 175344300 12394.260 *
##    3) Age_08_04< 32.5 93 817453300 18229.280
##      6) HP< 113 84 486744800 17703.850
##        12) Age_08_04>=21 36 113705000 16050.970 *
##        13) Age_08_04< 21 48 200924400 18943.500 *
##      7) HP>=113 9 91070000 23133.330 *

```

```
rpart.plot(dtm, type = 3, digits = 3, fallen.leaves = TRUE)
```



i.


Age_08_04, KM, Automatic_airco and HP are the important car specifications for predicting the car's price in order of high importance.

ii)

```
printcp(dtm)

##
## Regression tree:
## rpart(formula = Price ~ Age_08_04 + KM + Fuel_Type_Petrol +
Fuel_Type_Diesel +
##      Fuel_Type_CNG + Fuel_Type_ + HP + Automatic + Doors + Quarterly_Tax +
##      Mfr_Guarantee + Guarantee_Period + Airco + Automatic_airco +
##      CD_Player + Powered_Windows + Sport_Model + Tow_Bar, data = train,
##      method = "anova")
##
## Variables actually used in tree construction:
## [1] Age_08_04      Fuel_Type_Diesel HP
##
## Root node error: 9357382922/716 = 13068971
##
## n=716 (5 observations deleted due to missingness)
##
##      CP nsplit rel error  xerror    xstd
## 1 0.649396      0  1.00000 1.00103 0.083422
## 2 0.116779      1  0.35060 0.36471 0.030831
## 3 0.025610      2  0.23382 0.24973 0.026524
## 4 0.021313      3  0.20822 0.24331 0.026409
## 5 0.018394      4  0.18690 0.22770 0.023467
## 6 0.017393      5  0.16851 0.22084 0.022731
## 7 0.011372      6  0.15112 0.20620 0.022193
## 8 0.010000      7  0.13974 0.20452 0.022644

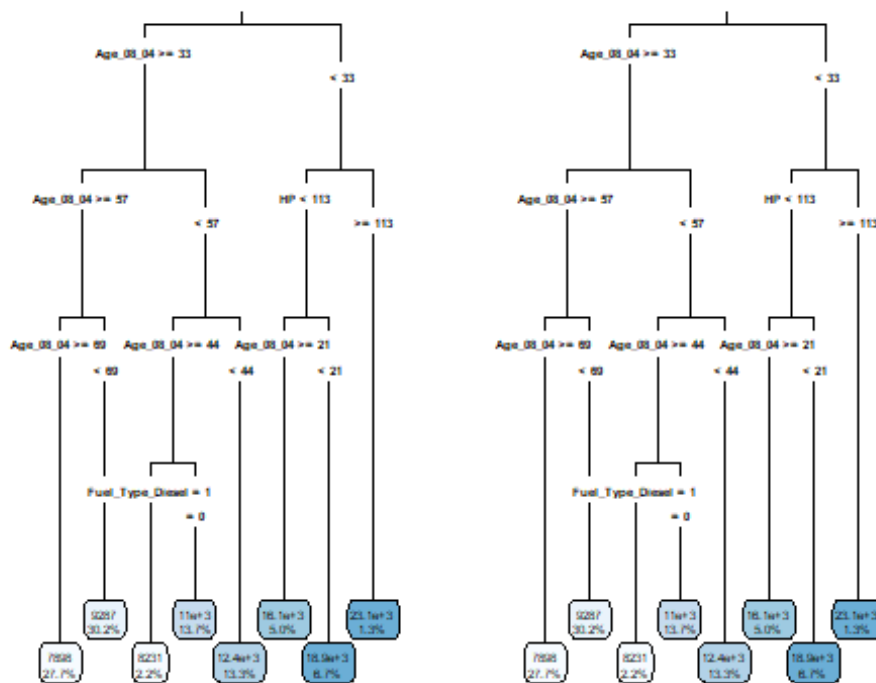
plotcp(dtm)
```



```
ptree<- prune(dtm, cp= dtm$cptable[which.min(dtm$cptable[, "xerror"]), "CP"])
```

comparing pruned tree to original decision tree

```
par(mfrow = c(1,2))
rpart.plot(dtm, type = 3, digits = 3, fallen.leaves = TRUE)
rpart.plot(ptree, type = 3, digits = 3, fallen.leaves = TRUE)
```



```
par(mfrow = c(1,1))
```

different predictions

```
p0 <- predict(dtm, train)
unique(p0)
```

```
## [1] 18943.500 12394.263 9287.106 11012.388 7897.672 23133.333 16050.972
## [8] 8230.938
```

```
p1 <- predict(dtm, validation)
unique(p1)
```

```
## [1] 12394.263 11012.388 7897.672 9287.106 18943.500 16050.972 8230.938
## [8] 23133.333
```

```
p2 <- predict(dtm, test)
p2
```

```
##          598          1029          821          430          690          309          604
## 11012.388  9287.106  9287.106 11012.388  9287.106 12394.263  9287.106
##          48          1427          413          1293          223          737          862
## 16050.972  7897.672  8230.938  7897.672 12394.263  9287.106  9287.106
##          889          1057          580          825          527          103          1141
##  9287.106  7897.672 11012.388  9287.106 11012.388 18943.500  7897.672
##          740          376          721          842          815          944          510
##  9287.106 12394.263  9287.106  9287.106  9287.106  9287.106 11012.388
##          996          847          491          1000          636          570          374
```

##	9287.106	9287.106	11012.388	9287.106	9287.106	11012.388	12394.263
##	741	969	449	504	34	178	909
##	9287.106	9287.106	11012.388	11012.388	16050.972	18943.500	9287.106
##	1411	945	293	106	421	735	377
##	7897.672	9287.106	12394.263	18943.500	11012.388	9287.106	12394.263
##	650	1281	253	1270	31	485	1105
##	9287.106	7897.672	12394.263	7897.672	16050.972	11012.388	7897.672
##	304	1138	409	1367	203	522	63
##	12394.263	7897.672	11012.388	7897.672	12394.263	11012.388	16050.972
##	1137	279	378	1371	25	470	659
##	7897.672	12394.263	12394.263	7897.672	16050.972	11012.388	9287.106
##	714	306	159	954	1197	427	56
##	9287.106	12394.263	18943.500	9287.106	7897.672	11012.388	16050.972
##	332	188	249	516	357	654	624
##	12394.263	12394.263	12394.263	11012.388	12394.263	9287.106	9287.106
##	287	628	674	185	207	477	812
##	12394.263	9287.106	9287.106	18943.500	12394.263	11012.388	9287.106
##	297	1083	514	1443	511	146	1
##	11012.388	7897.672	11012.388	9287.106	11012.388	18943.500	16050.972
##	827	584	285	743	802	1176	794
##	9287.106	11012.388	12394.263	9287.106	9287.106	7897.672	9287.106
##	870	277	1017	311	1091	1263	635
##	9287.106	12394.263	9287.106	12394.263	7897.672	7897.672	9287.106
##	333	948	792	150	37	349	894
##	12394.263	9287.106	9287.106	18943.500	16050.972	12394.263	9287.106
##	301	935	709	175	951	1296	251
##	12394.263	9287.106	9287.106	18943.500	9287.106	7897.672	12394.263
##	526	1241	661	89	660	814	72
##	11012.388	7897.672	9287.106	18943.500	9287.106	9287.106	16050.972
##	98	1268	931	274	1352	455	852
##	18943.500	7897.672	9287.106	12394.263	7897.672	11012.388	9287.106
##	574	1318	833	643	985	410	157
##	11012.388	7897.672	9287.106	9287.106	9287.106	11012.388	18943.500
##	572	30	1388	962	386	1429	1085
##	11012.388	16050.972	7897.672	9287.106	11012.388	7897.672	7897.672
##	1055	435	1332	52	594	190	671
##	7897.672	11012.388	7897.672	16050.972	11012.388	12394.263	9287.106
##	610	459	474	422	1274	1015	148
##	9287.106	8230.938	11012.388	11012.388	7897.672	9287.106	18943.500
##	963	1250	118	201	848	589	708
##	9287.106	7897.672	18943.500	11012.388	9287.106	11012.388	9287.106
##	1217	566	388	105	1026	1252	17
##	7897.672	11012.388	11012.388	18943.500	9287.106	7897.672	23133.333
##	919	978	1375	705	123	227	282
##	9287.106	9287.106	7897.672	9287.106	18943.500	12394.263	12394.263
##	551	1199	260	1299	7	181	918
##	11012.388	7897.672	12394.263	7897.672	16050.972	18943.500	9287.106
##	1305	846	838	1121	513	257	1159
##	7897.672	9287.106	9287.106	7897.672	11012.388	12394.263	7897.672
##	479	348	1403	362	1331	278	1109

```
## 11012.388 12394.263 7897.672 12394.263 7897.672 12394.263 7897.672
##      878      197      925      876      1209      1081      507
## 9287.106 8230.938 9287.106 9287.106 7897.672 7897.672 11012.388
##      545      1409      1051      626      1238      229      863
## 11012.388 7897.672 7897.672 9287.106 7897.672 12394.263 9287.106
##      620      868      286      218      585      161      658
## 9287.106 9287.106 12394.263 12394.263 11012.388 18943.500 9287.106
##      1113      1231      1423      1394      463      1087      392
## 7897.672 7897.672 7897.672 7897.672 11012.388 7897.672 8230.938
##      1136      937      761      321      341      82      219
## 7897.672 9287.106 9287.106 12394.263 12394.263 16050.972 11012.388
##      104      329      1205      826      1272      121      530
## 18943.500 12394.263 7897.672 9287.106 7897.672 18943.500 11012.388
##      1264      790      330      49      1380      965      198
## 7897.672 9287.106 12394.263 16050.972 7897.672 9287.106 12394.263
##      1157      666      989      804      1097      209      994
## 7897.672 9287.106 9287.106 9287.106 7897.672 12394.263 9287.106
##      691      983      328      371      830      1148      981
## 9287.106 9287.106 12394.263 12394.263 9287.106 7897.672 9287.106
##      550      50      1075      436      20      1251      433
## 11012.388 23133.333 7897.672 11012.388 16050.972 7897.672 11012.388
##      988      1179      991      805      845      599      756
## 9287.106 7897.672 9287.106 9287.106 9287.106 11012.388 9287.106
##      437
## 11012.388
```

RMSE Values for train data

```
difference_train = p0 - train$Price
diff_train<-difference_train^2
which(is.na(diff_train))

## 1443 1439 1437 1438 1441
##    66  359  436  607  678

diff_train = replace(diff_train, which(is.na(diff_train)), 0)
which(is.na(diff_train))

## named integer(0)

d_train<-mean(diff_train)
rmse_train<- sqrt(d_train)
rmse_train

## [1] 1346.716
```

RMSE Values for validation data

```
difference_valid = p1 - validation$Price
diff_valid<-difference_valid^2
which(is.na(diff_valid))

## 1443 1442 1439
##    73    77   274

diff_valid = replace(diff_valid, which(is.na(diff_valid)), 0)
which(is.na(diff_valid))

## named integer(0)

d_valid<-mean(diff_valid)
rmse_valid<- sqrt(d_valid)
rmse_valid

## [1] 1347.012
```

RMSE Values for test data

```
difference_test = p2 - test$Price
diff_test<-difference_test^2
which(is.na(diff_test))

## 1443
##    95


diff_test = replace(diff_test, which(is.na(diff_test)), 0)
which(is.na(diff_test))

## named integer(0)

d_test<-mean(diff_test)
rmse_test<- sqrt(d_test)
rmse_test

## [1] 1276.153

par(mfrow=c(1,3))
boxplot(p0)
boxplot(p1)
boxplot(p2)
```

```
par(mfrow = c(1,1))
```

Looking at the boxplots we observe our model performs good on test data; has a smaller spread as compared to the other two. There are more outliers in test because training tries to captures as many relationship as it can without overfitting.

classification tree

filling in the null values

```
summary(cars$Price)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##   4350    8450    9900   10731   11950   32500         7

cars$Price = ifelse(is.na(cars$Price),
                    ave(cars$Price, FUN = function(x) mean(x,
na.rm = TRUE))),
                    cars$Price)
```

creating new variable binned_price

```
library(Hmisc)

## Loading required package: lattice
```

```
## Loading required package: survival

##
## Attaching package: 'survival'

## The following object is masked from 'package:rpart':
##
##      solder

## Loading required package: Formula

## Loading required package: ggplot2

##
## Attaching package: 'Hmisc'

## The following objects are masked from 'package:base':
##
##      format.pval, units

cars$binprice <- as.numeric(cut2(cars$Price, g=20))
```

creating a new split of data

```
sample_train_ct<- sample(seq_len(nrow(cars)), size = floor(0.50*nrow(cars)))
sample_valid_ct<- sample(seq_len(nrow(cars)), size = floor(0.30*nrow(cars)))
sample_test_ct <- sample(seq_len(nrow(cars)), size = floor(0.20*nrow(cars)))

train_ct      <- cars[sample_train_ct, ]
validation_ct<- cars[sample_valid_ct, ]
test_ct       <- cars[sample_test_ct, ]
```

b.

Developing classification tree

```
classtree<-rpart(binprice ~
Age_08_04+KM+Fuel_Type_Petrol+Fuel_Type_Diesel+Fuel_Type_CNG+Fuel_Type_HP+
Automatic+Doors+Quarterly_Tax+Mfr_Guarantee+
Guarantee_Period+Airco+Automatic_airco+CD_Player+
Powered_Windows+Sport_Model+Tow_Bar , data = train_ct ,method = "class")

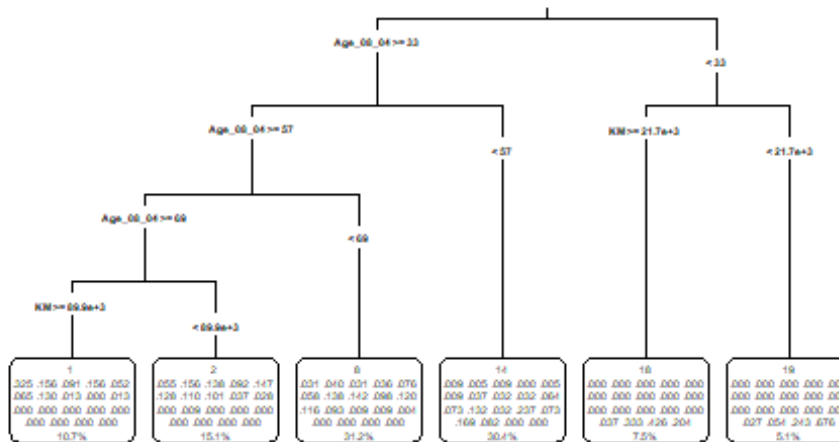
classtree

## n= 721
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
##  1) root 721 660 7 (0.055 0.054 0.043 0.042 0.053 0.047 0.085 0.071 0.046
0.062 0.058 0.071 0.012 0.075 0.024 0.055 0.053 0.044 0.05)
```

```
##      2) Age_08_04>=32.5 630 569 7 (0.063 0.062 0.049 0.048 0.06 0.054 0.097
0.081 0.052 0.071 0.067 0.081 0.014 0.086 0.027 0.059 0.029 0 0)
##      4) Age_08_04>=56.5 411 358 7 (0.092 0.092 0.071 0.073 0.09 0.078 0.13
0.11 0.063 0.075 0.063 0.054 0.0049 0.0049 0.0024 0 0 0 0)
##      8) Age_08_04>=68.5 186 155 1 (0.17 0.16 0.12 0.12 0.11 0.1 0.12
0.065 0.022 0.022 0 0.0054 0 0 0 0 0)
##     16) KM>=89869.5 77  52 1 (0.32 0.16 0.091 0.16 0.052 0.065 0.13
0.013 0 0.013 0 0 0 0 0 0 0 0) *
##     17) KM< 89869.5 109  92 2 (0.055 0.16 0.14 0.092 0.15 0.13 0.11
0.1 0.037 0.028 0 0.0092 0 0 0 0 0 0) *
##      9) Age_08_04< 68.5 225 193 8 (0.031 0.04 0.031 0.036 0.076 0.058
0.14 0.14 0.098 0.12 0.12 0.093 0.0089 0.0089 0.0044 0 0 0 0) *
##      5) Age_08_04< 56.5 219 167 14 (0.0091 0.0046 0.0091 0 0.0046 0.0091
0.037 0.032 0.032 0.064 0.073 0.13 0.032 0.24 0.073 0.17 0.082 0 0) *
##      3) Age_08_04< 32.5 91  55 19 (0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0.033 0.22
0.35 0.4)
##      6) KM>=21700 54  31 18 (0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0.037 0.33 0.43
0.2) *
##      7) KM< 21700 37  12 19 (0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0.027 0.054
0.24 0.68) *
```

```
rpart.plot(classtree, type = 3, digits = 3, fallen.leaves = TRUE)
```

```
## Warning: All boxes will be white (the box.palette argument will be
ignored) because
## the number of classes predicted by the model 19 is greater than
length(box.palette) 6.
## To silence this warning use box.palette=0 or trace=-1.
```




```

printcp(classtree)

##
## Classification tree:
## rpart(formula = binned_price ~ Age_08_04 + KM + Fuel_Type_Petrol +
##   Fuel_Type_Diesel + Fuel_Type_CNG + Fuel_Type_ + HP + Automatic +
##   Doors + Quarterly_Tax + Mfr_Guarantee + Guarantee_Period +
##   Airco + Automatic_airco + CD_Player + Powered_Windows + Sport_Model +
##   Tow_Bar, data = train_ct, method = "class")
##
## Variables actually used in tree construction:
## [1] Age_08_04 KM
##
## Root node error: 660/721 = 0.9154
##
## n= 721
##
##      CP nsplit rel error  xerror   xstd
## 1 0.060606      0  1.00000 1.00000 0.011322
## 2 0.018182      2  0.87879 0.88182 0.016049
## 3 0.015909      3  0.86061 0.91515 0.015000
## 4 0.010000      5  0.82879 0.89697 0.015594

plotcp(classtree)

```



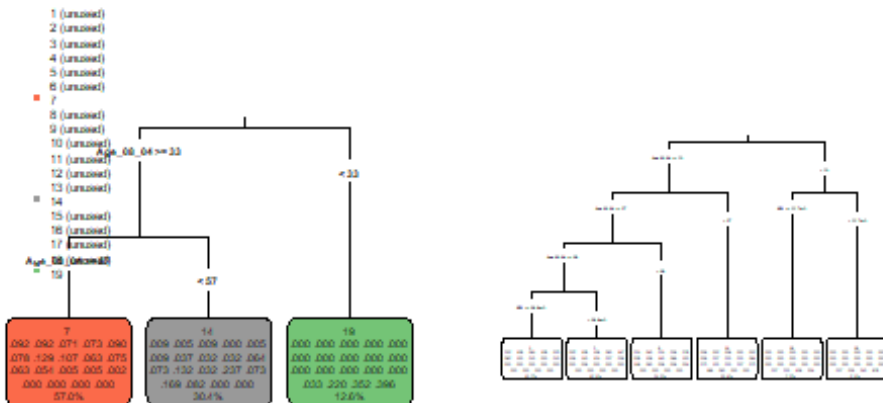
pruning tree

```
ptree1<- prune(classtree, cp=
classtree$cptable[which.min(classtree$cptable[, "xerror"]), "CP"])
ptree1

## n= 721
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 721 660 7 (0.055 0.054 0.043 0.042 0.053 0.047 0.085 0.071 0.046
0.062 0.058 0.071 0.012 0.075 0.024 0.055 0.053 0.044 0.05)
##   2) Age_08_04>=32.5 630 569 7 (0.063 0.062 0.049 0.048 0.06 0.054 0.097
0.081 0.052 0.071 0.067 0.081 0.014 0.086 0.027 0.059 0.029 0 0)
##     4) Age_08_04>=56.5 411 358 7 (0.092 0.092 0.071 0.073 0.09 0.078 0.13
0.11 0.063 0.075 0.063 0.054 0.0049 0.0049 0.0024 0 0 0 0) *
##     5) Age_08_04< 56.5 219 167 14 (0.0091 0.0046 0.0091 0 0.0046 0.0091
0.037 0.032 0.032 0.064 0.073 0.13 0.032 0.24 0.073 0.17 0.082 0 0) *
##   3) Age_08_04< 32.5 91  55 19 (0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0.033 0.22
0.35 0.4) *
```

```
par(mfrow= c(1,2))
rpart.plot(ptree1, type = 3, digits = 3, fallen.leaves = TRUE)
rpart.plot(classtree, type = 3, digits = 3, fallen.leaves = TRUE)
```

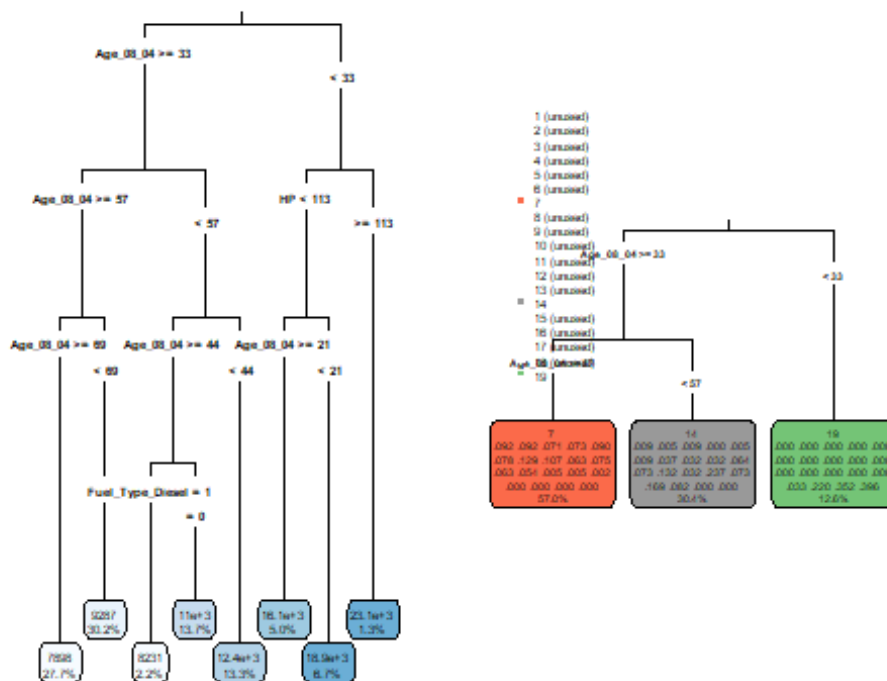
```
## Warning: All boxes will be white (the box.palette argument will be
## ignored) because
## the number of classes predicted by the model 19 is greater than
length(box.palette) 6.
## To silence this warning use box.palette=0 or trace=-1.
```



```
par(mfrow= c(1,1))
```

comparing regression tree and classification tree

```
par(mfrow= c(1,2))
rpart.plot(ptree, type = 3, digits = 3, fallen.leaves = TRUE)
rpart.plot(ptree1, type = 3, digits = 3, fallen.leaves = TRUE)
```



```
par(mfrow= c(1,1))
```

i.

The tree generated for CT and RT are different. Variable importance for CT is Age_08_04, KM, CD_Player & Quarterly_Tax while for RT it is Age_08_04, Automatic_airco, Quarterly_tax and HP. Size of trees are also differing. As we convert the price into bins, the variation which was present due to price being a continuous variable no longer exists and the variation of predictor variable is measured against the binned values due to which the effect reduces and less important variables for RT might be influential in CT.

```
p0_ct <- predict(classtree, train_ct)
unique(p0_ct)
```

```
##           1           2           3           4           5           6
## 155  0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
## 710  0.03111111 0.04000000 0.03111111 0.03555556 0.07555556 0.05777778
## 279  0.00913242 0.00456621 0.00913242 0.00000000 0.00456621 0.00913242
## 45   0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
## 1132 0.32467532 0.15584416 0.09090909 0.15584416 0.05194805 0.06493506
## 1223 0.05504587 0.15596330 0.13761468 0.09174312 0.14678899 0.12844037
##           7           8           9          10          11          12
## 155  0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
## 710  0.13777778 0.14222222 0.09777778 0.12000000 0.11555556 0.09333333
## 279  0.03652968 0.03196347 0.03196347 0.06392694 0.07305936 0.13242091
## 45   0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
## 1132 0.12987013 0.01298701 0.00000000 0.01298701 0.00000000 0.00000000
```

```
## 1223 0.11009174 0.10091743 0.03669725 0.02752294 0.00000000 0.009174312
##           13           14           15           16           17           18
## 155 0.000000000 0.000000000 0.000000000 0.02702703 0.05405405 0.2432432
## 710 0.008888889 0.008888889 0.004444444 0.00000000 0.00000000 0.0000000
## 279 0.031963470 0.237442922 0.073059361 0.16894977 0.08219178 0.0000000
## 45  0.000000000 0.000000000 0.000000000 0.03703704 0.33333333 0.4259259
## 1132 0.000000000 0.000000000 0.000000000 0.00000000 0.00000000 0.0000000
## 1223 0.000000000 0.000000000 0.000000000 0.00000000 0.00000000 0.0000000
##           19
## 155 0.6756757
## 710 0.0000000
## 279 0.0000000
## 45  0.2037037
## 1132 0.0000000
## 1223 0.0000000
```

```
p1_ct <- predict(classtree, validation_ct)
unique(p1_ct)
```

```
##           1           2           3           4           5           6
## 1083 0.32467532 0.15584416 0.09090909 0.15584416 0.05194805 0.06493506
## 862  0.03111111 0.04000000 0.03111111 0.03555556 0.07555556 0.05777778
## 118  0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.0000000
## 1396 0.05504587 0.15596330 0.13761468 0.09174312 0.14678899 0.12844037
## 17   0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.0000000
## 571  0.00913242 0.00456621 0.00913242 0.00000000 0.00456621 0.00913242
##           7           8           9           10          11          12
## 1083 0.12987013 0.01298701 0.00000000 0.01298701 0.00000000 0.00000000
## 862  0.13777778 0.14222222 0.09777778 0.12000000 0.11555556 0.09333333
## 118  0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.0000000
## 1396 0.11009174 0.10091743 0.03669725 0.02752294 0.00000000 0.009174312
## 17   0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.0000000
## 571  0.03652968 0.03196347 0.03196347 0.06392694 0.07305936 0.132420091
##           13          14          15          16          17          18
## 1083 0.000000000 0.000000000 0.000000000 0.00000000 0.00000000 0.0000000
## 862  0.008888889 0.008888889 0.004444444 0.00000000 0.00000000 0.0000000
## 118  0.000000000 0.000000000 0.000000000 0.02702703 0.05405405 0.2432432
## 1396 0.000000000 0.000000000 0.000000000 0.00000000 0.00000000 0.0000000
## 17   0.000000000 0.000000000 0.000000000 0.03703704 0.33333333 0.4259259
## 571  0.031963470 0.237442922 0.073059361 0.16894977 0.08219178 0.0000000
##           19
## 1083 0.0000000
## 862  0.0000000
## 118  0.6756757
## 1396 0.0000000
## 17   0.2037037
## 571  0.0000000
```

```
p2_ct <- predict(classtree, test_ct)
unique(p2_ct)
```



```
##           1           2           3           4           5           6
## 1305 0.05504587 0.15596330 0.13761468 0.09174312 0.14678899 0.12844037
## 239  0.00913242 0.00456621 0.00913242 0.00000000 0.00456621 0.00913242
## 1186 0.32467532 0.15584416 0.09090909 0.15584416 0.05194805 0.06493506
## 625  0.03111111 0.04000000 0.03111111 0.03555556 0.07555556 0.05777778
## 8    0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
## 103  0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
##           7           8           9           10          11          12
## 1305 0.11009174 0.10091743 0.03669725 0.02752294 0.00000000 0.009174312
## 239  0.03652968 0.03196347 0.03196347 0.06392694 0.07305936 0.132420091
## 1186 0.12987013 0.01298701 0.00000000 0.01298701 0.00000000 0.000000000
## 625  0.13777778 0.14222222 0.09777778 0.12000000 0.11555556 0.093333333
## 8    0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.000000000
## 103  0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.000000000
##           13          14          15          16          17          18
## 1305 0.000000000 0.000000000 0.000000000 0.00000000 0.00000000 0.00000000
## 239  0.031963470 0.237442922 0.073059361 0.16894977 0.08219178 0.00000000
## 1186 0.000000000 0.000000000 0.000000000 0.00000000 0.00000000 0.00000000
## 625  0.008888889 0.008888889 0.004444444 0.00000000 0.00000000 0.00000000
## 8    0.000000000 0.000000000 0.000000000 0.03703704 0.33333333 0.4259259
## 103  0.000000000 0.000000000 0.000000000 0.02702703 0.05405405 0.2432432
##           19
## 1305 0.00000000
## 239  0.00000000
## 1186 0.00000000
## 625  0.00000000
## 8    0.2037037
## 103  0.6756757
```

#RMSE Values for train data in classification trees

```
difference_train_ct = p0_ct - train_ct$binmed_price
```

#difference

```
diff_train_ct<-difference_train_ct^2
```

#diff_train

```
which(is.na(diff_train_ct))
```

```
## integer(0)
```

```
#diff_train_ct = replace(diff_train_ct, which(is.na(diff_train_ct)), 0)
```

```
#which(is.na(diff_train_ct))
```

```
d_train_ct<-mean(diff_train_ct)
```

#d_train

```
rmse_train_ct<- sqrt(d_train_ct)
```

```
rmse_train_ct
```

```
## [1] 11.09538
```

#RMSE Values for validation data in classification trees

```
difference_valid_ct = p1_ct - validation_ct$binmed_price
```

```
diff_valid_ct<-difference_valid_ct^2
```

```
which(is.na(diff_valid_ct))
```

```
## integer(0)

#diff_valid_ct = replace(diff_valid_ct, which(is.na(diff_valid_ct)), 0)
#which(is.na(diff_valid_ct))
d_valid_ct<-mean(diff_valid_ct)
rmse_valid_ct<- sqrt(d_valid_ct)
rmse_valid_ct

## [1] 11.02331

#RMSE Values for test data in classification trees
difference_test_ct = p2_ct - test_ct$binprice
diff_test_ct<-difference_test_ct^2
which(is.na(diff_test_ct))

## integer(0)

#diff_test_ct = replace(diff_test_ct, which(is.na(diff_test_ct)), 0)
#which(is.na(diff_test_ct))
d_test_ct<-mean(diff_test_ct)
rmse_test_ct<- sqrt(d_test_ct)
rmse_test_ct

## [1] 10.9677
```

- ii. Predict the price, using the RT and the CT, of a used Toyota Corolla with the specifications listed in Table below.

```
newcar<- data.frame("Age_08_04" = c(77), "KM" = c(117000), "Fuel_Type_Petrol"
= c(1), "HP" = c(110), "Automatic" = c(0), "Doors" = c(5), "Quarterly_Tax" =
c(100), "Mfr_Guarantee" = c(0), "Guarantee_Period" = c(3), "Airco" = c(1),
"Automatic_airco" = c(0), "CD_Player" = c(0), "Powered_Windows" = c(0),
"Sport_Model" = c(0), "Tow_Bar" = c(1), "Fuel_Type_" = c(0), "Fuel_Type_Diesel"
= c(0), "Fuel_Type_CNG" = c(1))
#View(newcar)

pred <-predict(dtm,newcar)
pred

##          1
## 7897.672

pred1 <-predict(ptree,newcar)
pred1

##          1
## 7897.672

pred_ct <- predict(classtree, newcar)
pred_ct

##          1          2          3          4          5          6          7
## 1 0.3246753 0.1558442 0.09090909 0.1558442 0.05194805 0.06493506 0.1298701
```

```
##           8 9           10 11 12 13 14 15 16 17 18 19
## 1 0.01298701 0 0.01298701 0 0 0 0 0 0 0 0 0 0

pred_ct_1 <- predict(ptree1, newcar)
pred_ct_1

##           1           2           3           4           5           6
## 1 0.09245742 0.09245742 0.07055961 0.0729927 0.09002433 0.07785888
##           7           8           9          10          11          12
## 1 0.1289538 0.107056 0.06326034 0.07542579 0.06326034 0.05352798
##           13          14          15 16 17 18 19
## 1 0.00486618 0.00486618 0.00243309 0 0 0 0 0
```

Problem 2

```
Banks = read.csv("Banks.csv")
summary(Banks)

##      Obs      Financial.Condition TotCap.Assets  TotExp.Assets
## Min.   : 1.00   Min.   :0.0           Min.   : 0.700   Min.   :0.0700
## 1st Qu.: 5.75   1st Qu.:0.0           1st Qu.: 7.125   1st Qu.:0.0800
## Median :10.50   Median :0.5           Median : 9.000   Median :0.1000
## Mean   :10.50   Mean   :0.5           Mean   : 9.395   Mean   :0.1035
## 3rd Qu.:15.25   3rd Qu.:1.0           3rd Qu.:12.325   3rd Qu.:0.1200
## Max.   :20.00   Max.   :1.0           Max.   :20.600   Max.   :0.1600
## TotLns.Lses.Assets
## Min.   :0.3000
## 1st Qu.:0.5350
## Median :0.6450
## Mean   :0.6325
## 3rd Qu.:0.7250
## Max.   :1.0400
```

Q2a

```
Banks$X1 <- ifelse(Banks$Financial.Condition ==1,"weak","strong")
Banks$X1<- factor(Banks$X1)
fit.full <- glm(X1 ~ TotCap.Assets+TotExp.Assets + TotLns.Lses.Assets,data=
Banks,family=binomial(link='logit'))

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

summary(fit.full)

##
## Call:
## glm(formula = X1 ~ TotCap.Assets + TotExp.Assets + TotLns.Lses.Assets,
##      family = binomial(link = "logit"), data = Banks)
##
## Deviance Residuals:
```

```

##           Min           1Q           Median           3Q           Max
## -3.464e-05 -2.100e-08  0.000e+00  2.100e-08  3.311e-05
##
## Coefficients:
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -622.53  470392.65  -0.001    0.999
## TotCap.Assets      -16.57   13113.35  -0.001    0.999
## TotExp.Assets    2535.76 2072975.47   0.001    0.999
## TotLns.Lses.Assets  774.36  595862.75   0.001    0.999
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2.7726e+01  on 19  degrees of freedom
## Residual deviance: 2.9026e-09  on 16  degrees of freedom
## AIC: 8
##
## Number of Fisher Scoring iterations: 25

fit.reduce <- fit.full <- glm(X1 ~ TotExp.Assets + TotLns.Lses.Assets,data=
Banks,family=binomial(link='logit'))
summary(fit.reduce)

##
## Call:
## glm(formula = X1 ~ TotExp.Assets + TotLns.Lses.Assets, family =
binomial(link = "logit"),
##      data = Banks)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.64035  -0.35514   0.02079   0.53234   1.03373
##
## Coefficients:
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -14.188     6.122  -2.317   0.0205 *
## TotExp.Assets     79.964    39.263   2.037   0.0417 *
## TotLns.Lses.Assets  9.173     6.864   1.336   0.1814
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 27.726  on 19  degrees of freedom
## Residual deviance: 12.831  on 17  degrees of freedom
## AIC: 18.831
##
## Number of Fisher Scoring iterations: 6

#Odds
coef(fit.reduce)

```

```
##      (Intercept)      TotExp.Assets TotLns.Lses.Assets
##      -14.187552         79.963941         9.173215

exp(coef(fit.reduce))

##      (Intercept)      TotExp.Assets TotLns.Lses.Assets
##      6.893258e-07      5.344393e+34      9.635549e+03

#Probability
Banks$predicted <-
predict(fit.reduce,newdata=subset(Banks,select=c(2,3,4,5)))
Banks$predicted_prob <-
predict(fit.reduce,newdata=subset(Banks,select=c(2,3,4,5)),type='response')
```

Q2b.

```
TotLns.Lses.Assets <- as.numeric(0.6)
TotExp.Assets <- as.numeric(0.11)
newbank <- data.frame(TotLns.Lses.Assets, TotExp.Assets)
#names(newbank)[1]<- "TotLns&Lses/Assets"
#names(newbank)[2]<- "TotExp/Assets"

#logit
NewLogit <- predict(fit.reduce,newbank)
NewLogit

##      1
## 0.1124105

#Probability
NewProbab <- predict(fit.reduce,newbank,type="response")
WeakStrong <- ifelse(NewProbab >= 0.5,1,0)
WeakStrong

## 1
## 1

#odds
odds<-NewProbab/(1-NewProbab)
odds

##      1
## 1.118972
```

The New Bank is classified under “Weak”.

Q2c.

```
Cutoff <- as.numeric(0.5)

#odds being financially weak
```

```
Odds <- Cutoff/(1-Cutoff)
Odds
```

```
## [1] 1
```

```
NewLogit <- log(Odds)
NewLogit
```

```
## [1] 0
```

Q2d.

```
coef(fit.reduce)
```

```
##      (Intercept)      TotExp.Assets TotLns.Lses.Assets
##      -14.187552         79.963941         9.173215
```

```
exp(coef(fit.reduce))
```

```
##      (Intercept)      TotExp.Assets TotLns.Lses.Assets
##      6.893258e-07      5.344393e+34      9.635549e+03
```

Total loans and & leases to assests for odds of being financially weak increase by 9.635549e+03.

Q2e. When a bank that is in poor financial condition is misclassified as financially strong, the misclassification cost is much higher than when a financially strong bank is misclassified as weak. To minimize the expected cost of misclassification, should the cutoff value for classification (which is currently at 0.5) be increased or decreased?

```
Banks$final <- ifelse(Banks$predicted_prob>=0.5,"weak","strong")
table(Banks$X1,Banks$final)
```

```
##
##      strong weak
## strong      9   1
## weak       0  10
```

```
Banks$final <- ifelse(Banks$predicted_prob>=0.7,"weak","strong")
table(Banks$X1,Banks$final)
```

```
##
##      strong weak
## strong      9   1
## weak       2   8
```

Ans: The Cutoff value should be decreased.

Problem 3

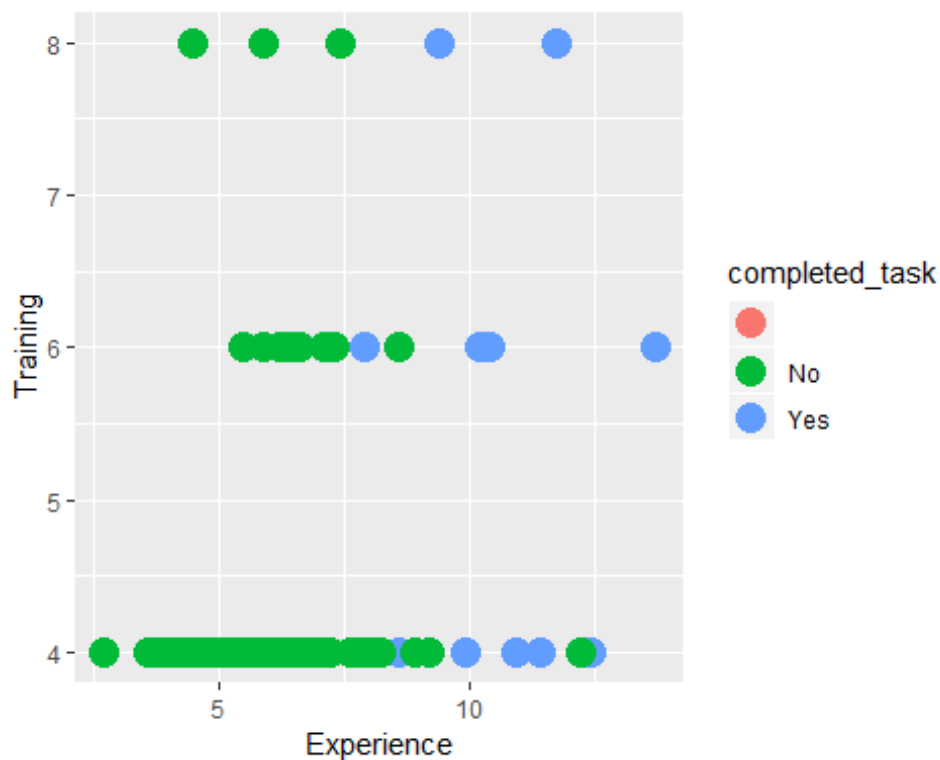
a

```
System_Administrators <- read.csv("System Administrators.csv")

names(System_Administrators)[3]<-"completed_task"
str(System_Administrators)

## 'data.frame':    80 obs. of  3 variables:
## $ Experience     : num  10.9 9.9 10.4 13.7 9.4 12.4 7.9 8.9 10.2 11.4 ...
## $ Training       : int   4 4 6 6 8 4 6 4 6 4 ...
## $ completed_task: Factor w/ 3 levels "", "No", "Yes": 3 3 3 3 3 3 3 3 3 3
ggplot(System_Administrators,aes(x=Experience,y=Training,color=completed_task
)) +
  geom_point(size=5)

## Warning: Removed 5 rows containing missing values (geom_point).
```



Experience is an important predictor which will help in classification of administrators into yes or no for completion of task. As we can see, administrators with more experience tend to complete the task. Training is being represented as not a strong predictor for classification

b.

```
System_Administrators$completed_task <-
factor(System_Administrators$completed_task)
str(System_Administrators)

## 'data.frame':    80 obs. of  3 variables:
## $ Experience     : num  10.9 9.9 10.4 13.7 9.4 12.4 7.9 8.9 10.2 11.4 ...
## $ Training       : int   4 4 6 6 8 4 6 4 6 4 ...
```

```
## $ completed_task: Factor w/ 3 levels "", "No", "Yes": 3 3 3 3 3 3 3 3 3 3
...
```

b.

```
model_logistic <- glm(completed_task ~
Experience+Training,data=System_Administrators,family=binomial(link='logit'))
System_Administrators$predicted <-
predict(model_logistic,System_Administrators[,c("Experience","Training")],typ
e="response")
System_Administrators$predicted_output <-
ifelse(System_Administrators$predicted>0.5,"Yes","No")
table(System_Administrators$completed_task,System_Administrators$predicted_ou
tput)
```

```
##
##      No Yes
##      0  0
## No   58  2
## Yes  5 10
```

33.33% of the programmers are incorrectly classified as failing to complete the task from among the programmers who complete the task.

c.

```
System_Administrators$predicted_output <-
ifelse(System_Administrators$predicted>0.5,"Yes","No")
table(System_Administrators$completed_task,System_Administrators$predicted_ou
tput)
```

```
##
##      No Yes
##      0  0
## No   58  2
## Yes  5 10
```

```
System_Administrators$predicted_output <-
ifelse(System_Administrators$predicted>0.6,"Yes","No")
table(System_Administrators$completed_task,System_Administrators$predicted_ou
tput)
```

```
##
##      No Yes
##      0  0
## No   59  1
## Yes  6  9
```

```
System_Administrators$predicted_output <-
ifelse(System_Administrators$predicted>0.4,"Yes","No")
table(System_Administrators$completed_task,System_Administrators$predicted_ou
tput)
```

```
##
##      No Yes
```



```
##      0    0
##   No  56    4
##   Yes  4   11
```

Ans: To decrease the percentage in part B we need to decrease the cutoff value as shown above.

d.

```
exp <- (log(1.01)+10.9813-0.1805*4)/1.1269
exp
## [1] 9.112832
```

If a programmer has experience more than 9.11 years with 4years of training than his probability of completing the task exceeds 50%